

The Constrained Application Protocol (CoAP)

Moosa Yahyazadeh

The University of Iowa

November 4, 2016

What is CoAP?


- The Constrained Application Protocol (CoAP) is a specialized web transfer protocol for use with constrained nodes and constrained (e.g., low-power, lossy) networks
- The protocol is designed for machine-to-machine (M2M) applications
 - smart energy
 - building automation
- It provides a request/response interaction model between application endpoints
 - One design goal → keep message overhead small
 - Why? limiting the need for fragmentation in constrained environments

Request/response interaction model

- Interaction model of CoAP is similar to the client/server model of HTTP
- Machine-to-Machine interaction result in CoAP implementation acting in both client and server roles
- CoAP request/response
 - a request is sent by a client for an action (using a Method Code) on a resource (identified by a URI) on a server
 - server then sends a response with a Response Code; this response may include a resource representation
- Unlike HTTP, CoAP deals with these interchanges asynchronously over a datagram-oriented transport such as UDP.
 - using messages layer that supports optional reliability

CoAP layers

- Abstract Layering of CoAP



- One could think of CoAP logically as using a two-layer approach, a CoAP messaging layer used to deal with UDP and the asynchronous nature of the interactions, and the request/response interactions using Method and Response Codes
- CoAP is however a single protocol, with messaging and request/response as just features of the CoAP header.

Messaging model

- Message types
 - Confirmable
 - Non-confirmable
 - Acknowledgement
 - Reset
- Method Codes and Response Codes included in some of these messages make them carry requests or responses
- The basic exchanges of the four types of messages are somewhat orthogonal to the request/response interactions; requests can be carried in Confirmable and Non-confirmable messages, and responses can be carried in these as well as piggybacked in Acknowledgement messages
 - Thus, Requests cannot be carried in Ack messages

Message format

- CoAP is based on the exchange of compact messages
 - transported over UDP → each CoAP message occupies the data section of one UDP datagram
 - may also be used over Datagram Transport Layer Security (DTLS)
 - It could also be used over other transports such as SMS, TCP, or SCTP
 - UDP-lite [RFC3828] and UDP zero checksum [RFC6936] are not supported by CoAP

Message format (Cont...)

- messages are encoded in a simple binary format
- message structure
 - header
 - token (if any)
 - options (if any)
 - payload marker (if there is any payload)
 - payload (if any)

Ver	T	TKL	Code	Message ID
Token (if any, TKL bytes) ...				
Options (if any) ...				
11111111			Payload (if any) ...	

Message format - Header

- header (fixed-size 4-byte)
 - Version (Ver): 2-bit unsigned integer
 - CoAP version number
 - Implementations of this specification **MUST** set this field to 1 (01 binary). Other values are reserved for future versions. Messages with unknown version numbers **MUST** be silently ignored
 - Type (T): 2-bit unsigned integer
 - Indicates if this message is of type Confirmable (0), Non-confirmable (1), Acknowledgement (2), or Reset (3)
 - Token Length (TKL): 4-bit unsigned integer
 - Indicates the length of the variable-length Token field (0-8 bytes)
 - Lengths 9-15 are reserved, **MUST NOT** be sent, and **MUST** be processed as a message format error

Message format - Header (Cont...)

- header (fixed-size 4-byte) (Cont...)
 - Code: 8-bit unsigned integer
 - split into a 3-bit class (most significant bits) and a 5-bit detail (least significant bits)
 - documented as "c.dd". "c" is a digit from 0 to 7 for the 3-bit subfield and "dd" are two digits from 00 to 31 for the 5-bit subfield.
 - The class can indicate a request (0), a success response (2), a client error response (4), or a server error response (5). (All other class values are reserved.)
 - As a special case, Code 0.00 indicates an Empty message.
 - In case of a request, the Code field indicates the Request Method; in case of a response, a Response Code.
 - Message ID: 16-bit unsigned integer

Message format - Header (Cont...)

- header (fixed-size 4-byte) (Cont...)
 - Message ID: 16-bit unsigned integer
 - split into a 3-bit class (most significant bits) and a 5-bit detail (least significant bits)
 - documented as "c.dd". "c" is a digit from 0 to 7 for the 3-bit subfield and "dd" are two digits from 00 to 31 for the 5-bit subfield.
 - The class can indicate a request (0), a success response (2), a client error response (4), or a server error response (5). (All other class values are reserved.)
 - As a special case, Code 0.00 indicates an Empty message.
 - In case of a request, the Code field indicates the Request Method; in case of a response, a Response Code.

Rules for generating a Message ID and matching messages

Code registries

Request/Response Semantics